

Native Client: Safe, Portable Execution of Untrusted x86 Native Code

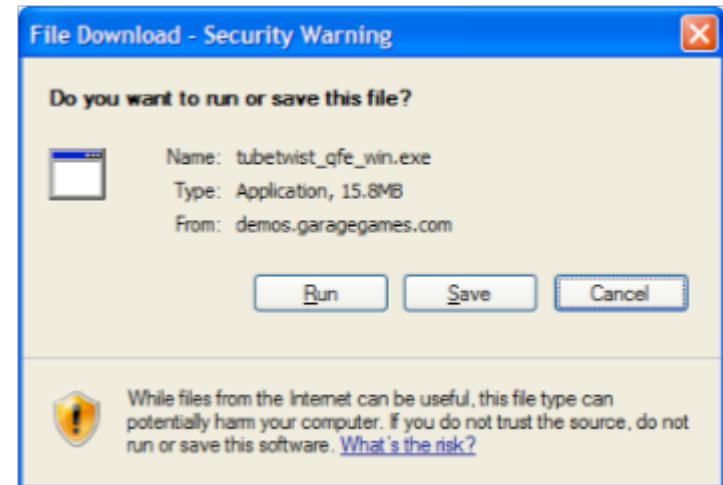


Brad Chen
January 2009

with Bennet Yee, David Sehr,
Shiki Okasaka, Robert Muth,
Greg Dardyk, Nicholas Fullagar,
Neha Narula, Tavis Ormandy

Safe Native Code for Web Apps

- Browser plugins deliver performance but sacrifice security
 - basic flaw: they require blind trust
 - exacerbated by weak OS security models



NativeClient enables native performance without sacrificing security

Native Client Goals

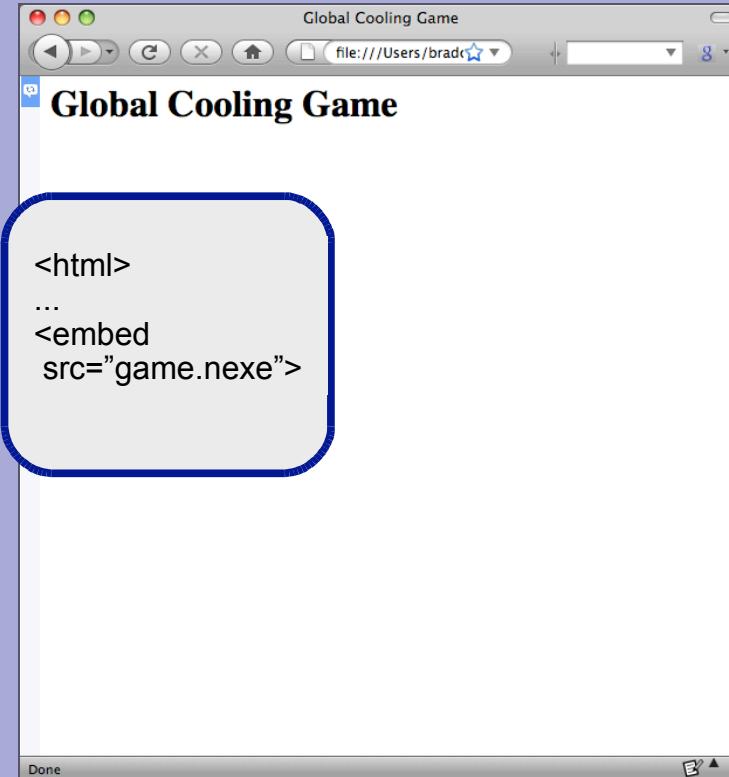
- ◆ Goals
 - Safe execution of untrusted x86-32 native code
 - Performance \approx native code
 - OS/browser portability
 - Trustworthy
 - Trusted code base is small and simple
 - Open source
 - Complement, not replace existing infrastructure
- ◆ Non-goals
 - execution of existing binary code (must recompile)
 - ISA portability

A Quick Demonstration

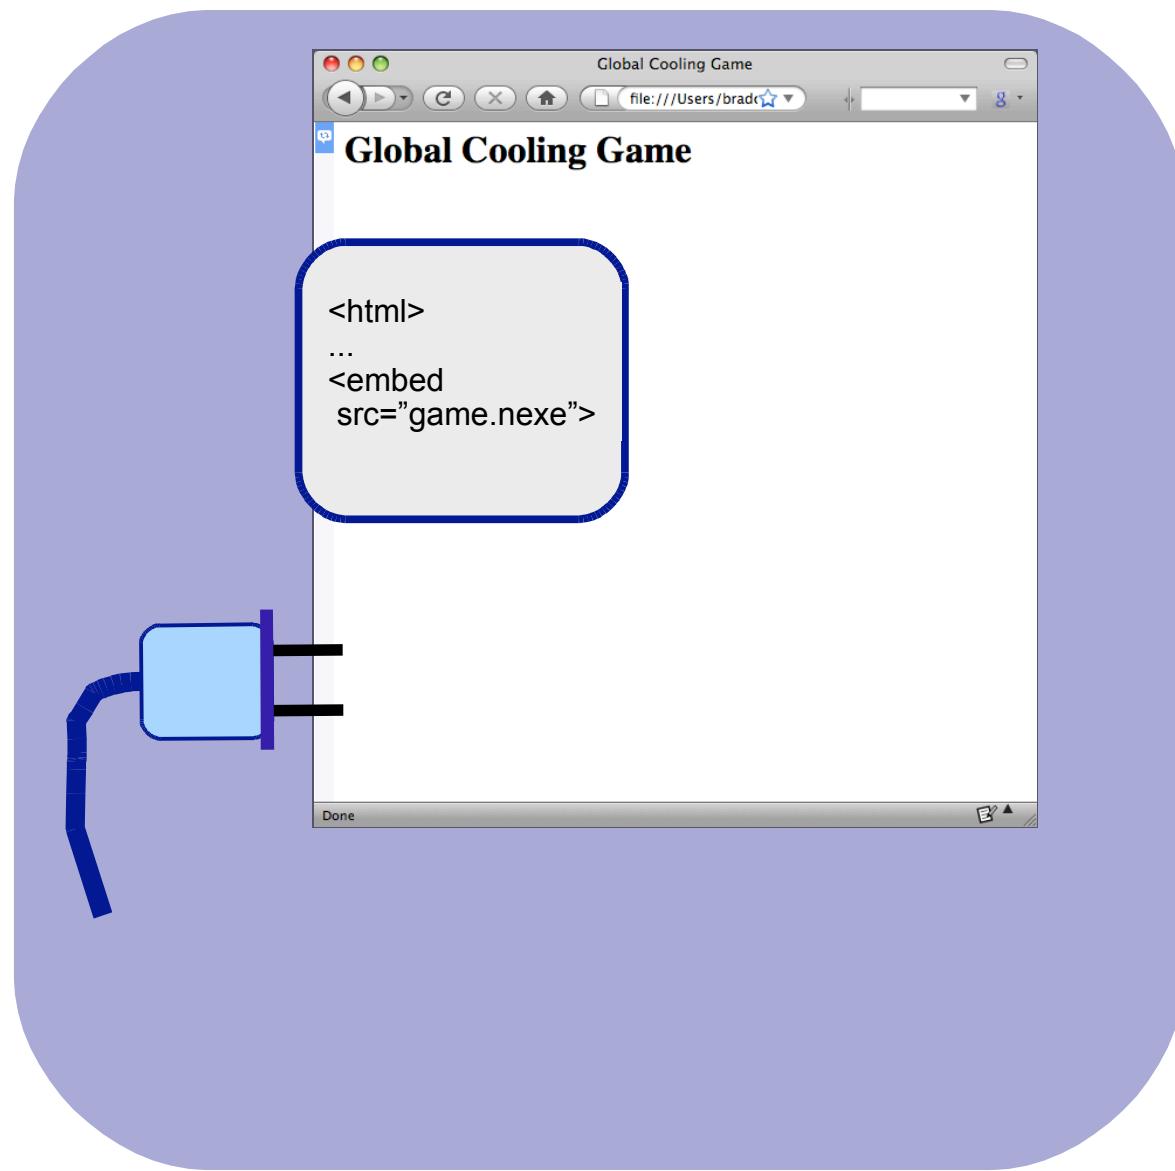
Life of a NaCl-enabled Web App



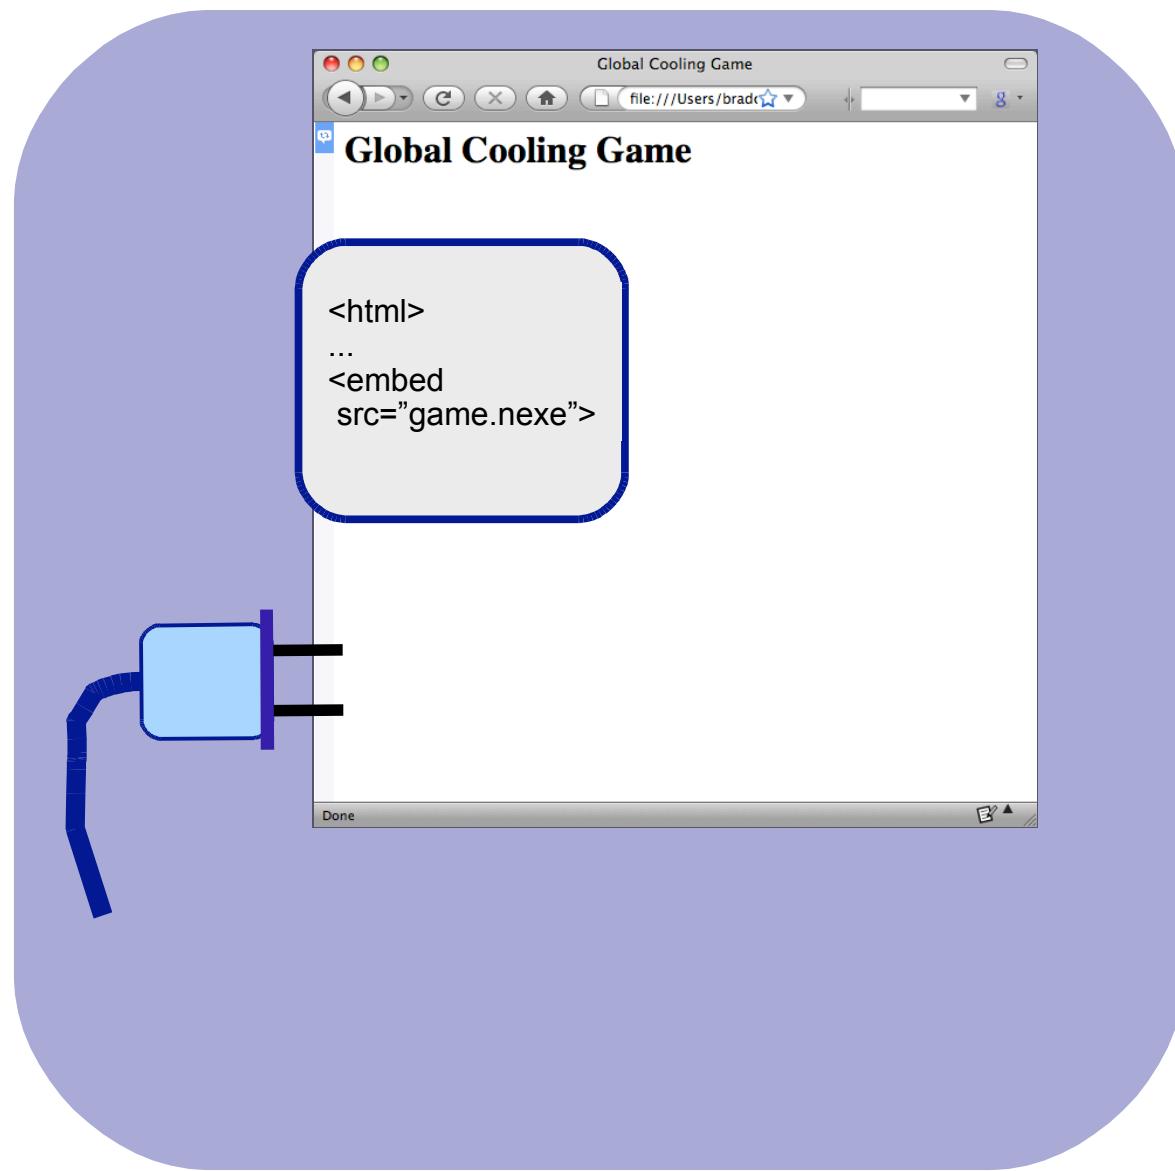
Life of a NaCl-enabled Web App



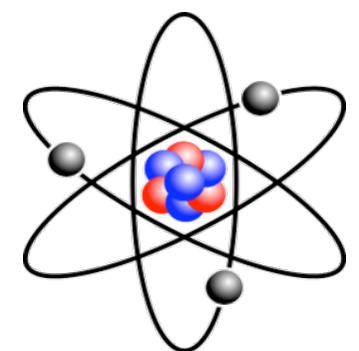
Life of a NaCl-enabled Web App



Life of a NaCl-enabled Web App

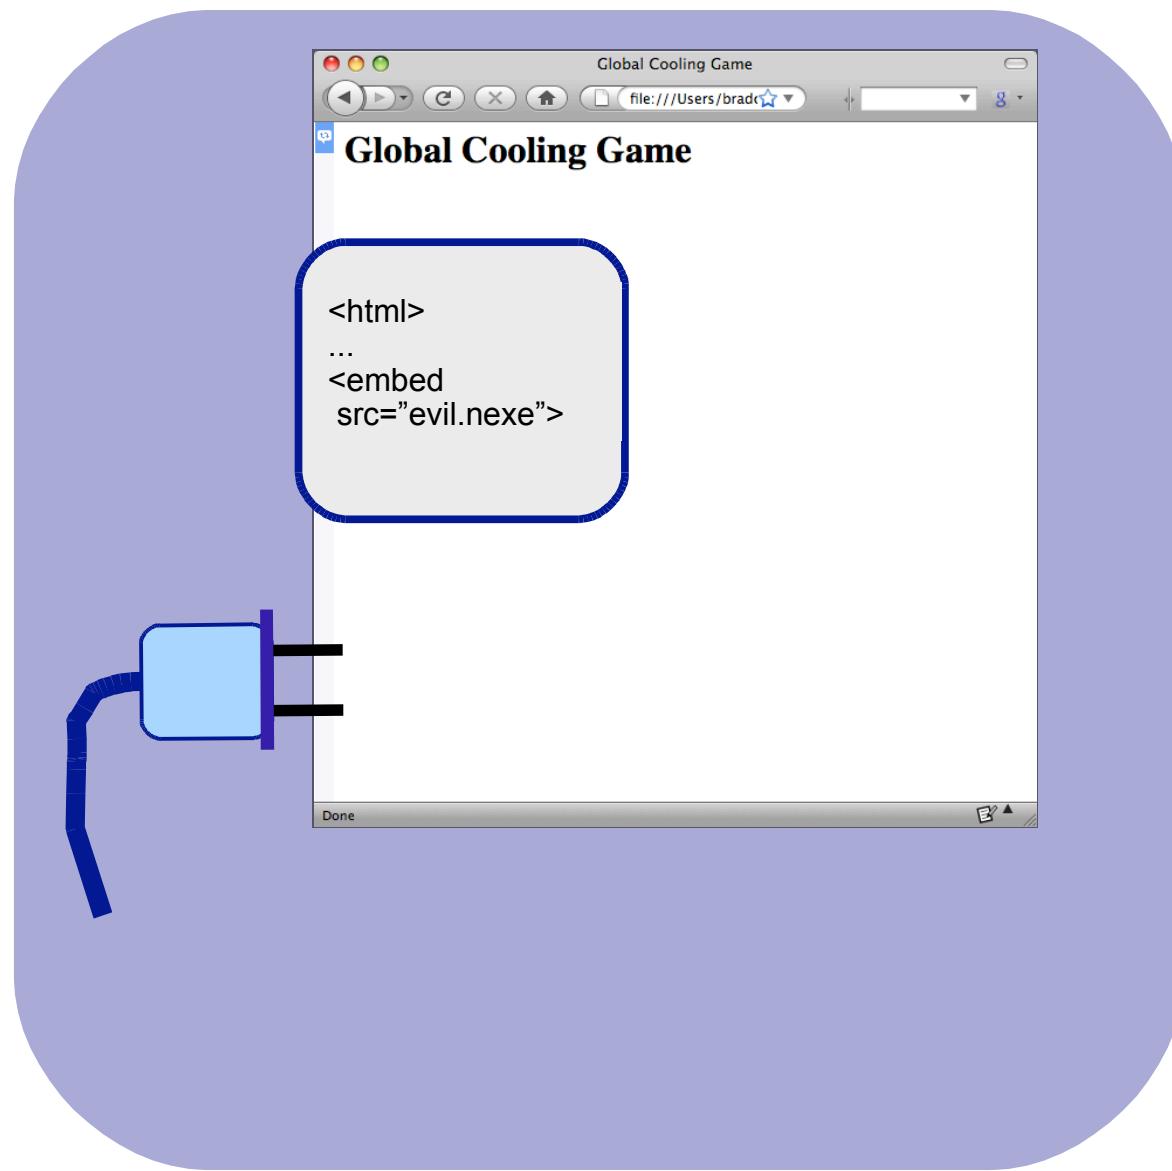


8



Google™

Life of a NaCl-enabled Web App

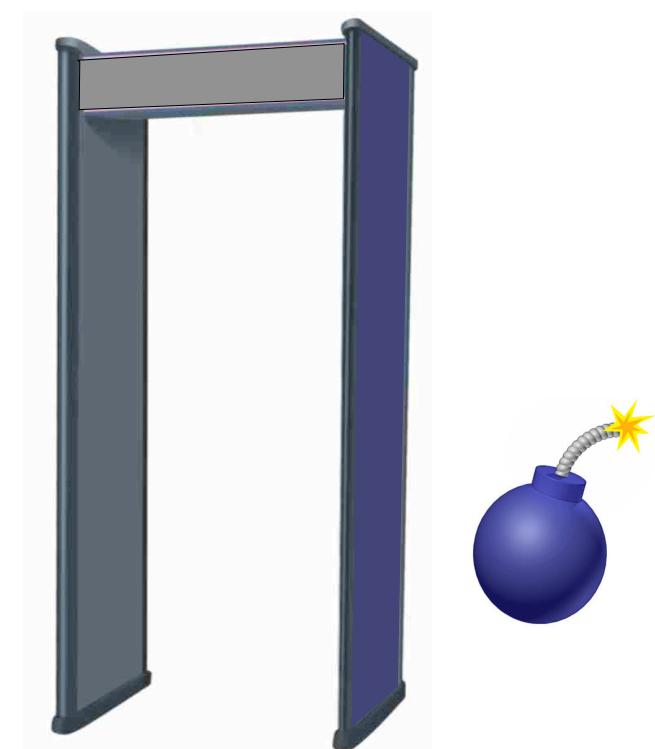
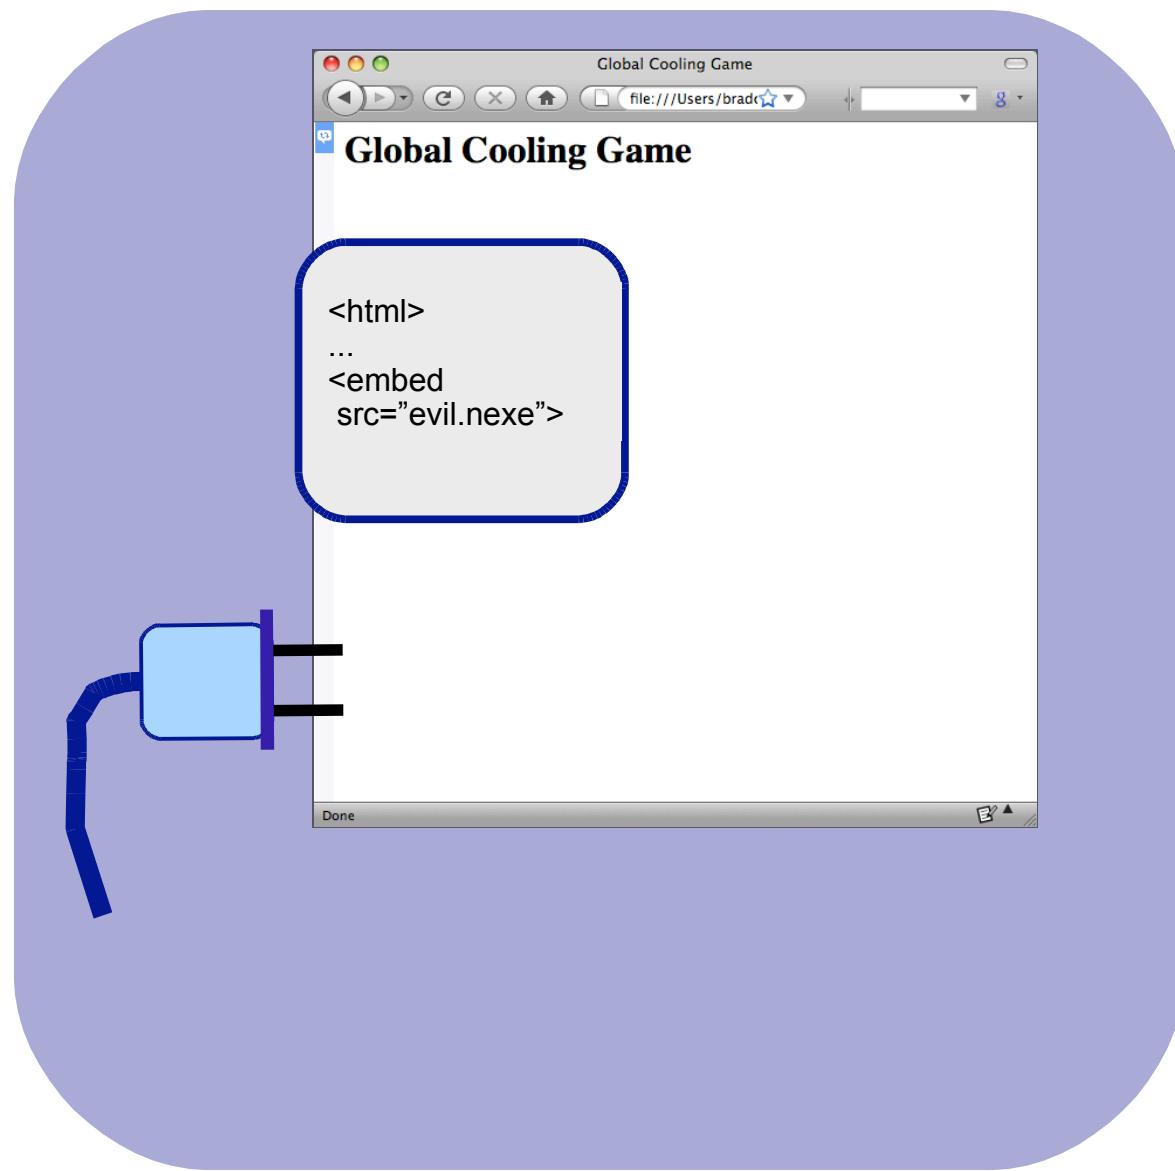


9

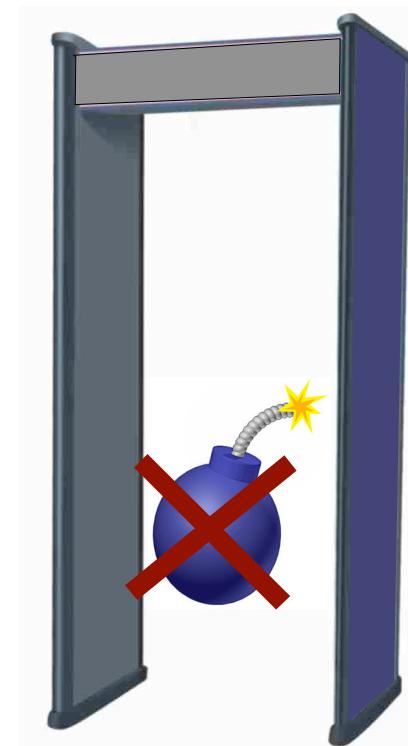
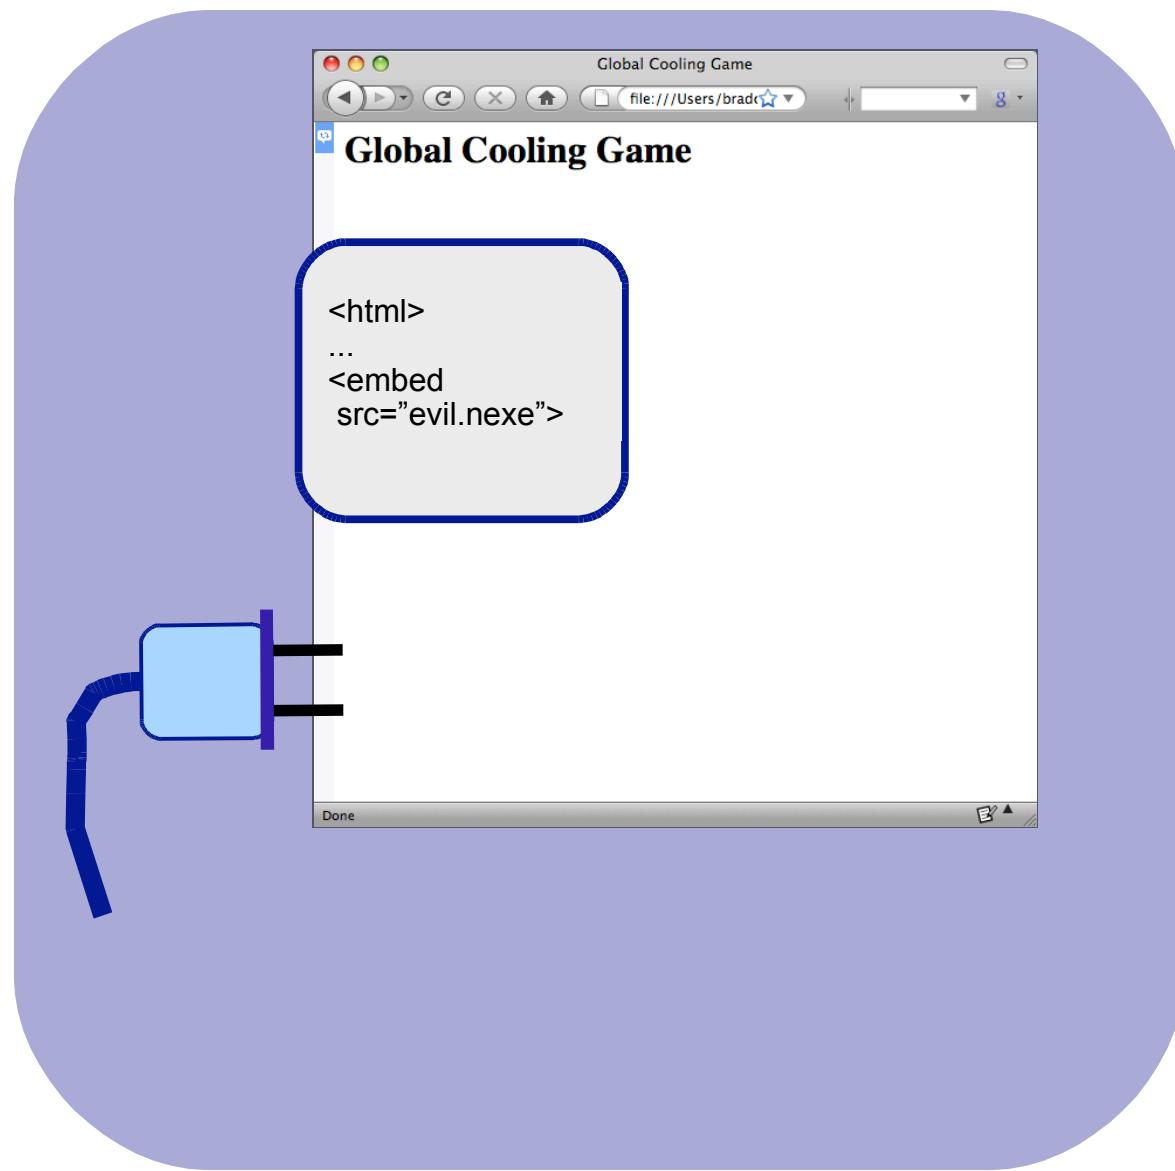


Google™

Life of a NaCl-enabled Web App

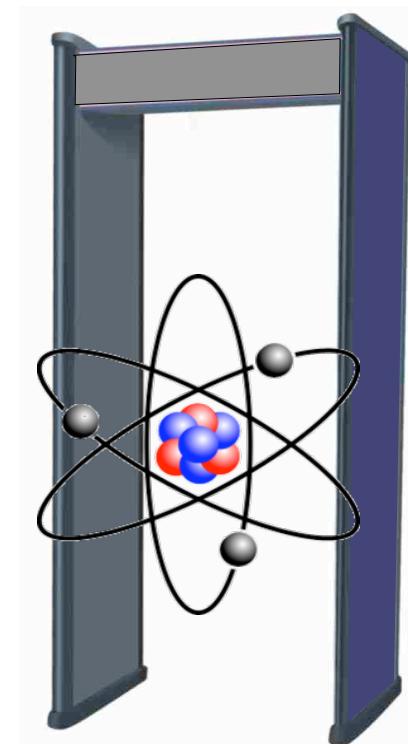
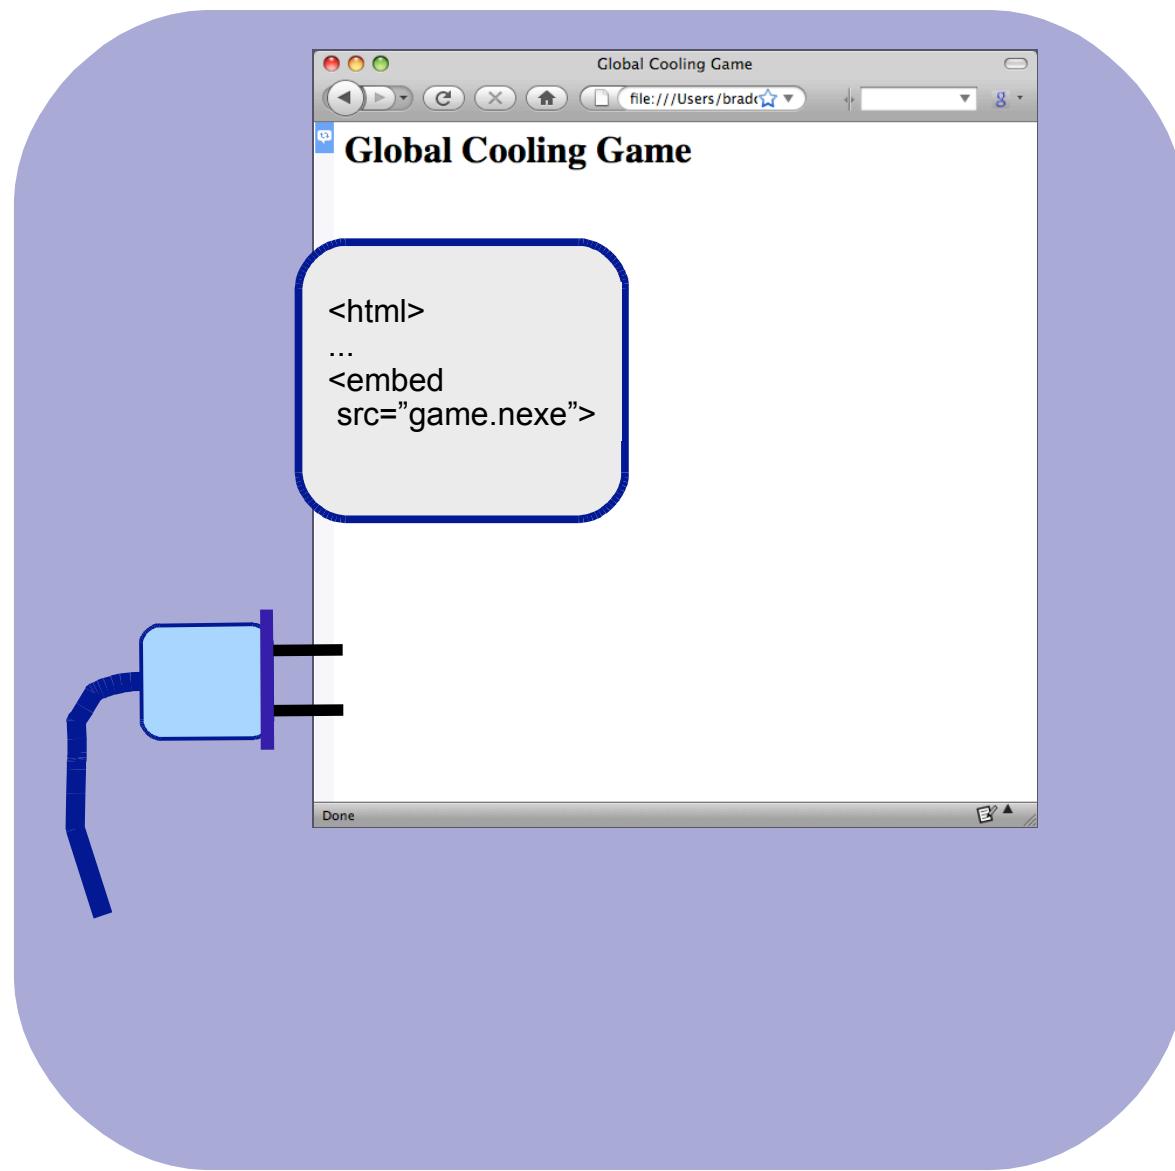


Life of a NaCl-enabled Web App



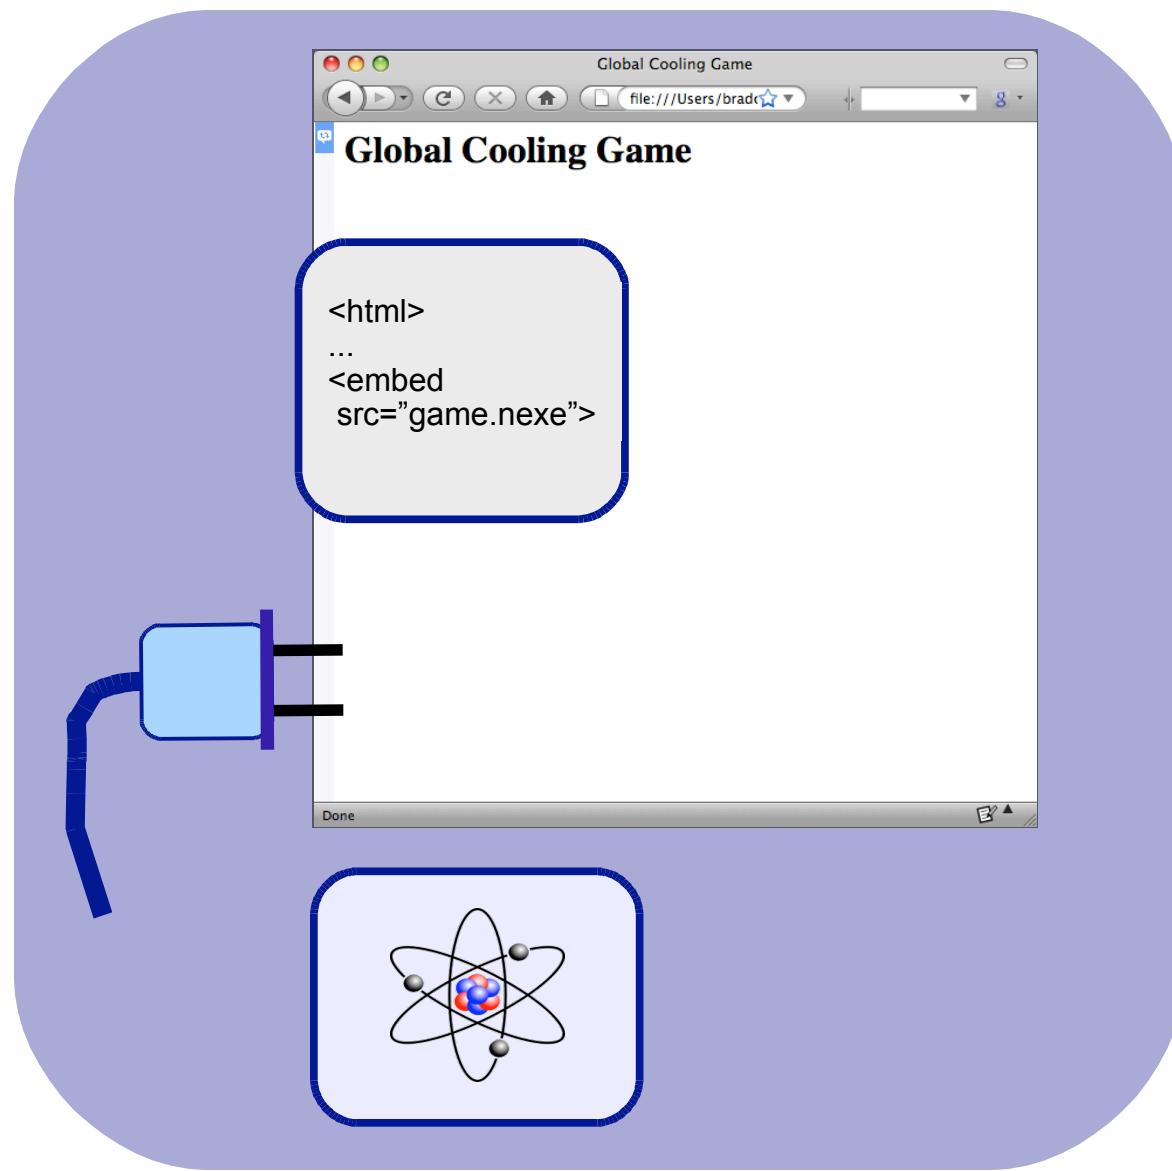
Google™

Life of a NaCl-enabled Web App

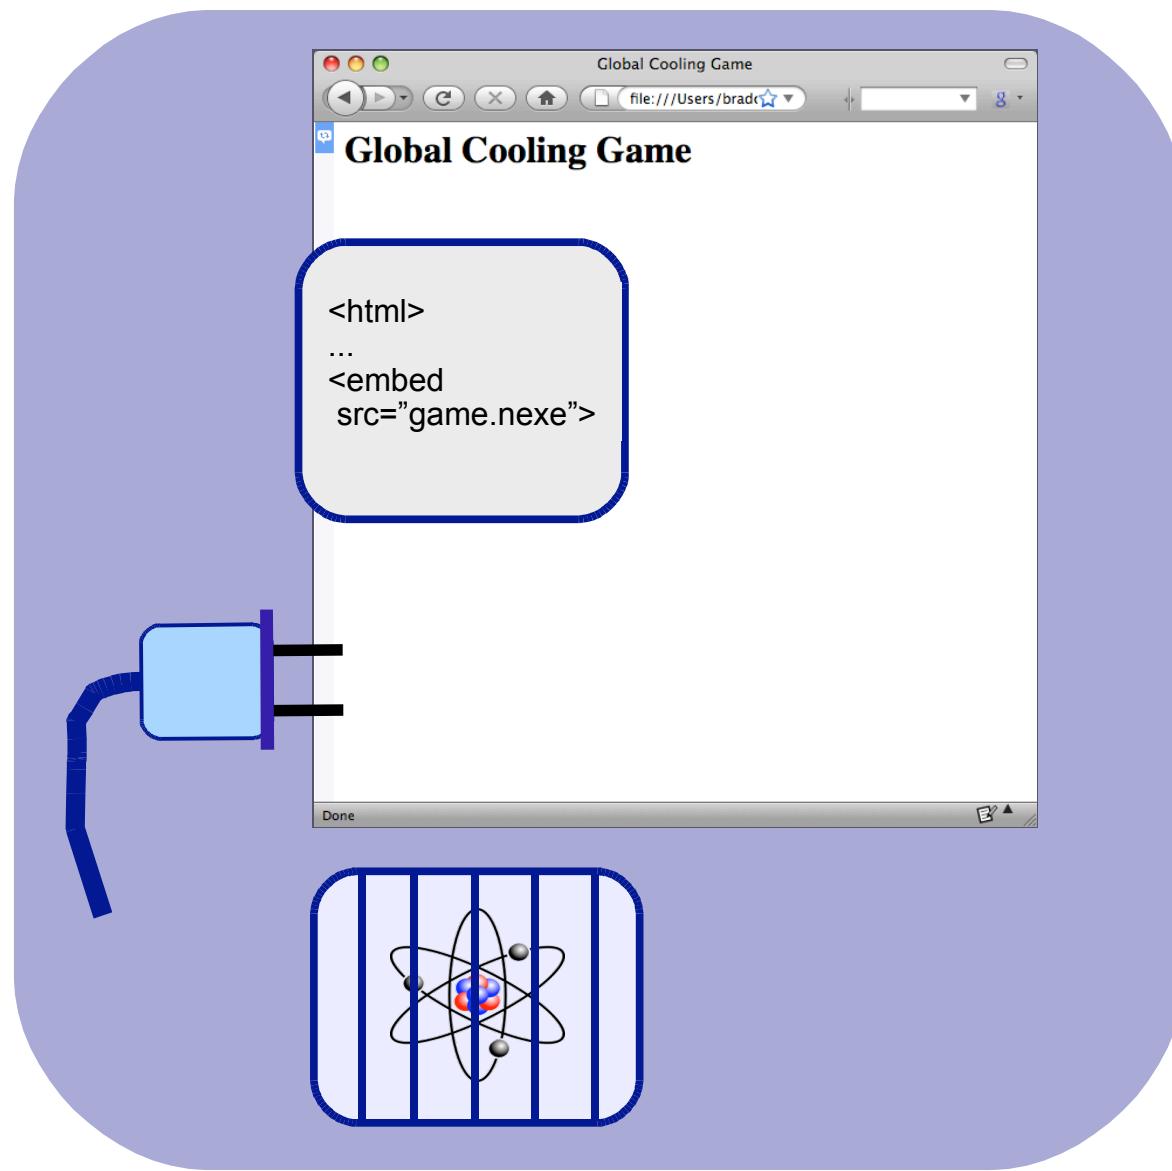


Google™

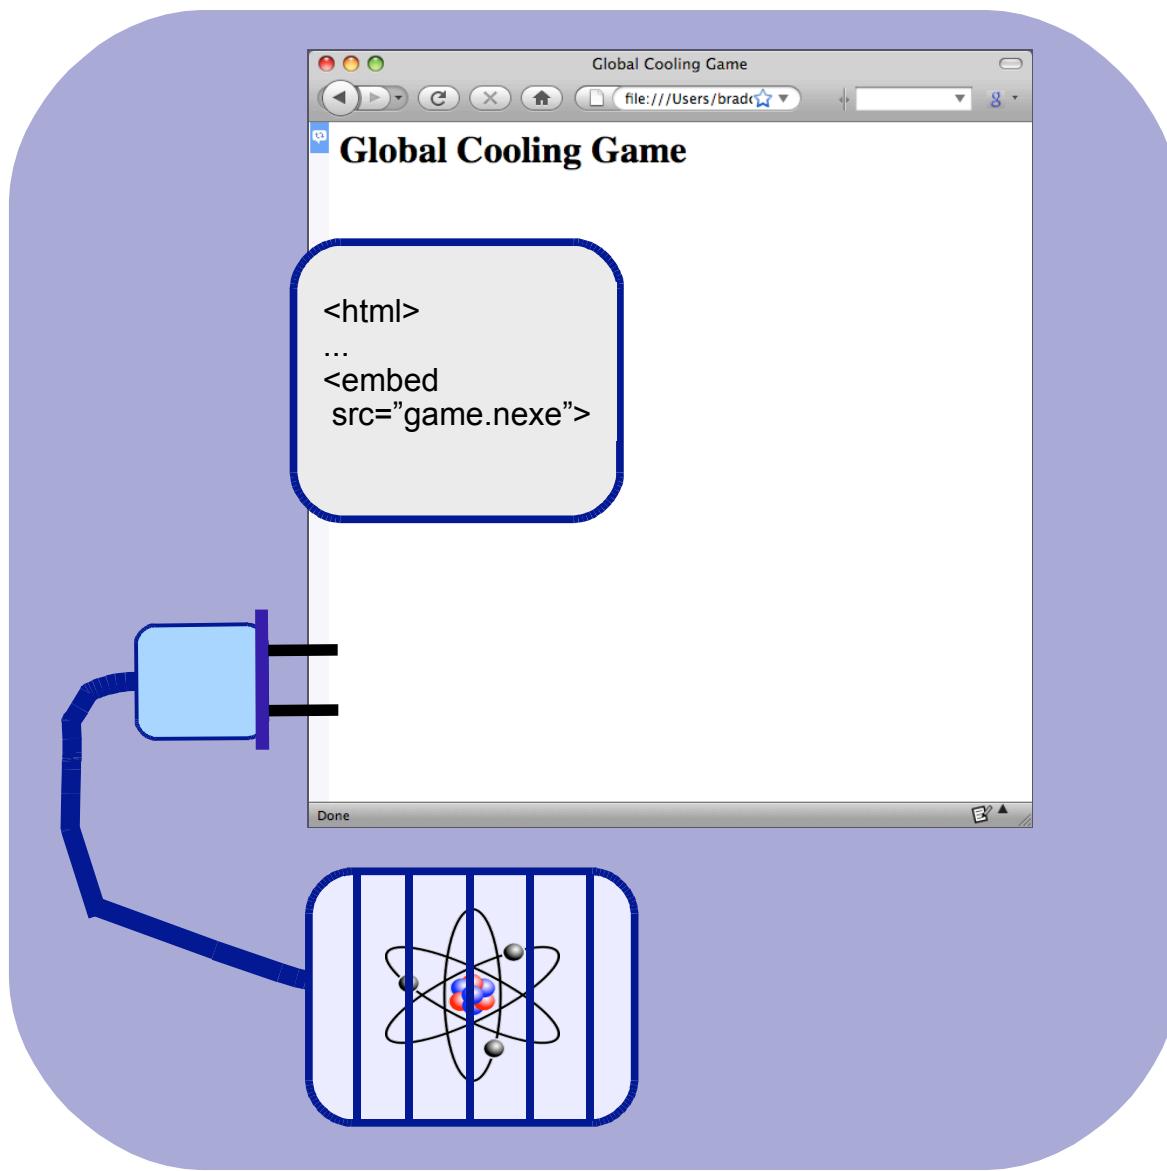
Life of a NaCl-enabled Web App



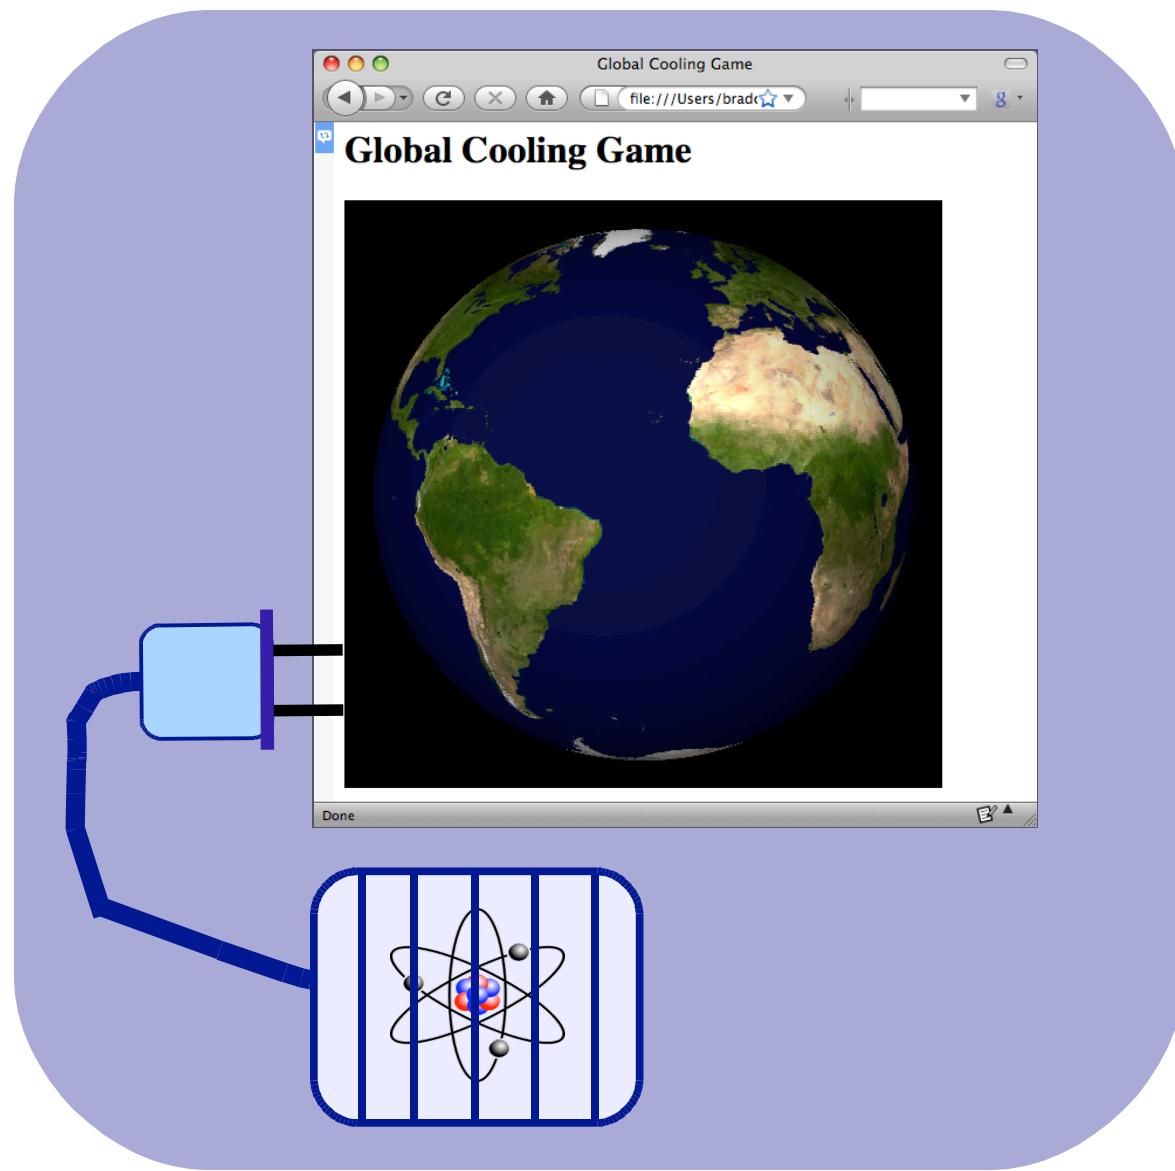
Life of a NaCl-enabled Web App



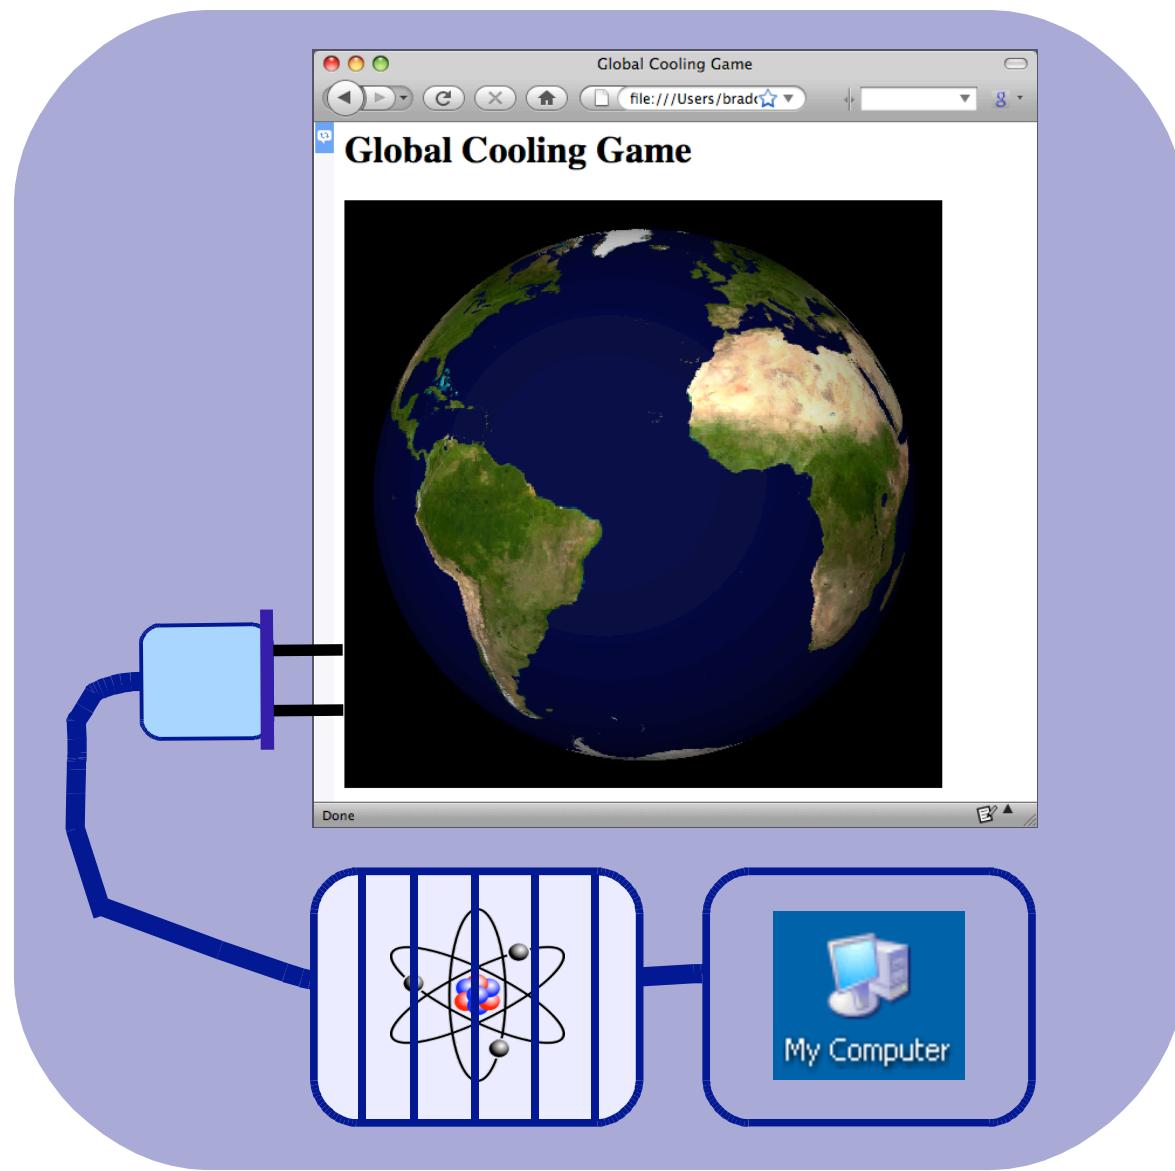
Life of a NaCl-enabled Web App



Life of a NaCl-enabled Web App



Life of a NaCl-enabled Web App



Q: What does “safe” mean?

Q: What does “safe” mean?

A: No side effects except via explicit secure interfaces.

Example 1: An Evil System Call

```
void DoEvil() {  
    char *eargv[] = {"/bin/rm", "-rf", "/home/*", NULL};  
    int rc = execve(eargv[0], eargv, NULL);  
}
```

NaCl modules are not allowed to directly execute system calls or other 'special' instructions.

Variants

- dynamically generated code
- code embedded in data
- overlapping instructions

Example 1a: An Evil System Call

```
// MacOS version
int TrustMe(int returnaddr1, const char *path,
            char *const argv[], char *const envp[]) {
    int immx = 0x0000340f;      // 0f 34 is syscall inst.
    int codeaddr = 14 + (int)TrustMe;

    asm("mov    $59, %eax");    // set syscall # for execve
    asm("add    $32, %esp");    // pop local storage
    asm("mov    %esp, %ecx");   // kernel wants esp in ecx
    asm("jmp    *-20(%ecx)");   // jump to syscall inst
                                // via pointer in codeaddr
}

#define NULL 0
char *const eargv[] = {
    "/bin/rm", "-rf", "/home/*", NULL};
int main(int argc, char *argv[]) {
    TrustMe(-1, eargv[0], eargv, NULL);
}
```

Example 1a: An Evil System Call

```
// MacOS version
int TrustMe(int returnaddr1, const char *path,
            char *const argv[], char *const envp[]) {
int immx = 0x0000340f; // 0f 34 is syscall inst.
int codeaddr = 14 + (int)TrustMe;

asm("mov    $59, %eax");      // set syscall # for execve
asm("add    $32, %esp");     // pop local storage
asm("mov    %esp, %ecx");    // kernel wants esp in ecx
asm("jmp    *-20(%ecx)"); // jump to syscall inst
                                // via pointer in codeaddr
}

#define NULL 0
char *const eargv[] = {
    "/bin/rm", "-rf", "/home/*", NULL};
int main(int argc, char *argv[]) {
    TrustMe(-1, eargv[0], eargv, NULL);
}
```

Example 2: A Wild Write

```
char* magicp = FindReturnAddrOnIEStack();  
*magicp = (char *) &DoEvil;
```

- NaCl modules should not be able to damage other computations on the system
- Assumption: operating system is not cooperative
 - Basic requirement for portability
 - Must constrain writes to known safe memory

Example 3: A Devious Read

```
char *creditcardptr = 0x38a390;  
GoShoppingWithFriends(creditcardptr);
```

- NaCl modules should only be allowed to read data inside their sandbox.

Catalog of Threats

Primary Threats

- Special instructions
- Wild control flow
- Wild writes
- Wild reads

Secondary Threats

- Infinite loops
- Memory leaks
- Corrupt output
- Covert channels

Catalog of Threats

Primary Threats

- Special instructions
 - Wild control flow
 - Wild writes
 - Wild reads
- Code Integrity
 - Control Flow Integrity
 - Data Integrity

Secondary Threats

- Infinite loops
- Memory leaks
- Corrupt output
- Covert channels

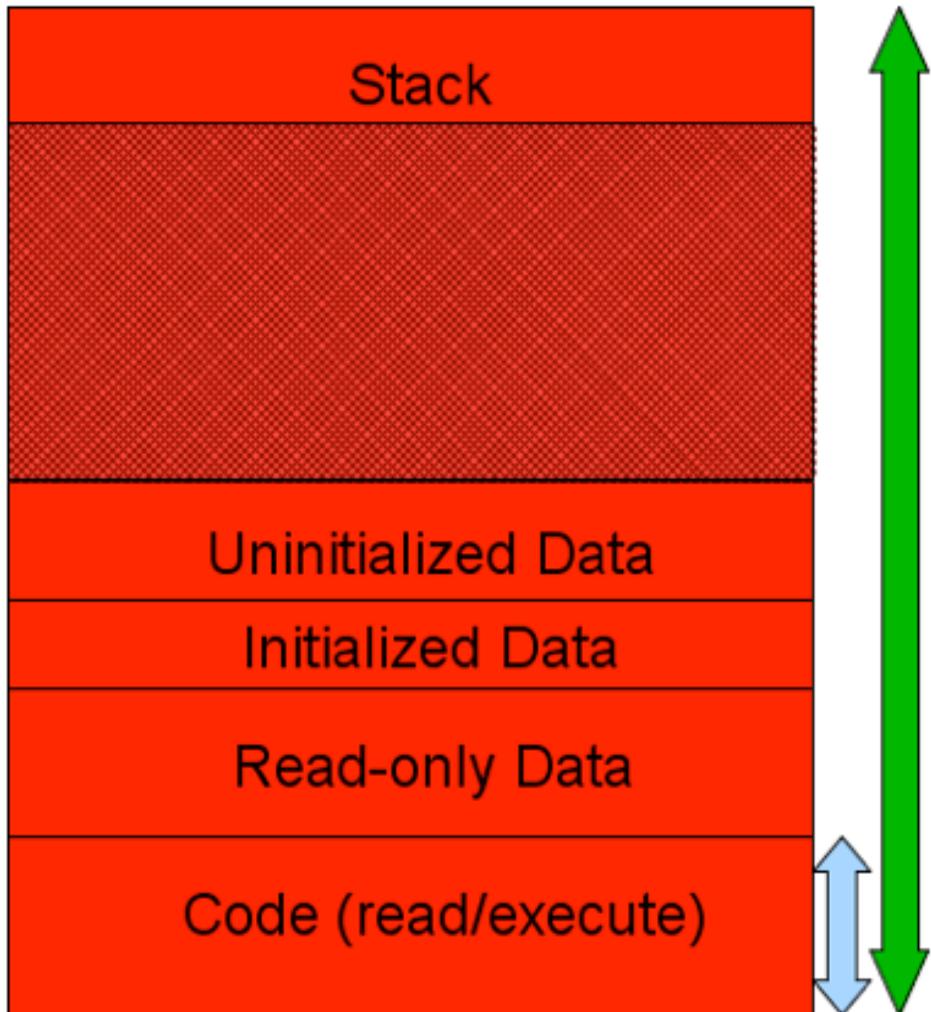
Outline

- Introduction
- Background
- Eliminating Side Effects
 - Inner Sandbox
 - Outer Sandbox
- Allowing Side-Effects Safely
- Evaluation

Inner Sandbox

- Software Fault Isolation + x86 segments
- No OS dependencies
 - portable
 - relatively immune to OS defects
 - caveat: LDT initialization
- Uses modified compiler, assembler and loader
 - modifications are simple and open
 - details follow
- Runtime overhead is very small

Data/Memory Integrity



NaCl's address space is bounded by x86 segments:

- **%cs** allows execution only in code segment
- **%ds**, **%es**, **%ss** restrict data address space
- **%gs** used for thread descriptor

NaCl Control Flow Integrity Rules

Making x86 disassembly reliable:

- All valid instruction boundaries are determined by a top-to-bottom disassembly of text segment
- All aligned memory addresses must be a valid instruction boundary
- Computed jumps use sandboxing to force alignment
- Branch targets are within code segment (dynamically enforced via segmentation)
- Forbidden: mixed text/data, overlapping instructions, dynamic code generation

Control Flow Integrity

- Static validation of most calls and jumps

```
call    8ef40
```

- Sandboxed returns and indirect calls/jumps

- NaCl text must be aligned to n -byte blocks
 - All indirect branch targets must be aligned

```
naclcall 3830(%eax)
```



```
lea    %eax, 3830(%eax)
and   %eax, 0xffffffffe0
call  %eax
```

NaCl ISA/ABI Restrictions

Given reliable disassembly, NaCl can enforce:

- Omitted instructions: `ret`, `int`, `syscall` ... most x86-64 omissions
- Restricted prefix usage
 - At most one prefix byte per instructions
 - Segment, addr16 prefixes disallowed
 - Prefix usage restricted to known useful prefixes
- However, ISA extensions generally supported
 - unsupported instructions overwritten HLT

The Most Recent Validator Bug

...

```
83 e2 e0    and 0xfffffe0, %eax  
0f e2        jmp (%eax)      ; ; okay
```

...

```
83 e2 e0    and 0xfffffe0, %eax  
0f 12        jmp (*%eax)    ; ; not okay
```

Outer Sandbox

Outer Sandbox

- System call filter for untrusted modules
- Redundant: “defense-in-depth”
- OS-dependent
- Implementation
 - Linux: ptrace
 - MacOS: sandbox.h
 - Windows: system call ACLs

Outline

- Introduction
- Background
- Eliminating Side Effects (SFI)
- Allowing Side-Effects Safely
- Evaluation

NaCl Runtime Environment

- IMC
 - reliable datagram service
 - handles: shared memory, descriptors
- POSIX subset
 - operations on descriptors
 - memory management
- pthreads, synchronization, TLS
- multimedia: video, audio, events

NaCl Runtime: Interesting features

- Shared memory for communication into/out of untrusted modules
- Calls out to trusted code via 'trampoline' and 'springboard' mechanism
- Uses a separate trusted stack
- All hardware exceptions are fatal
 - C++ catch/throw exceptions are okay
- Calls into untrusted code via simple RPC system, with constrained data types:
 - int float string array
 - opaque “handle”; no pointers

Attack Surface

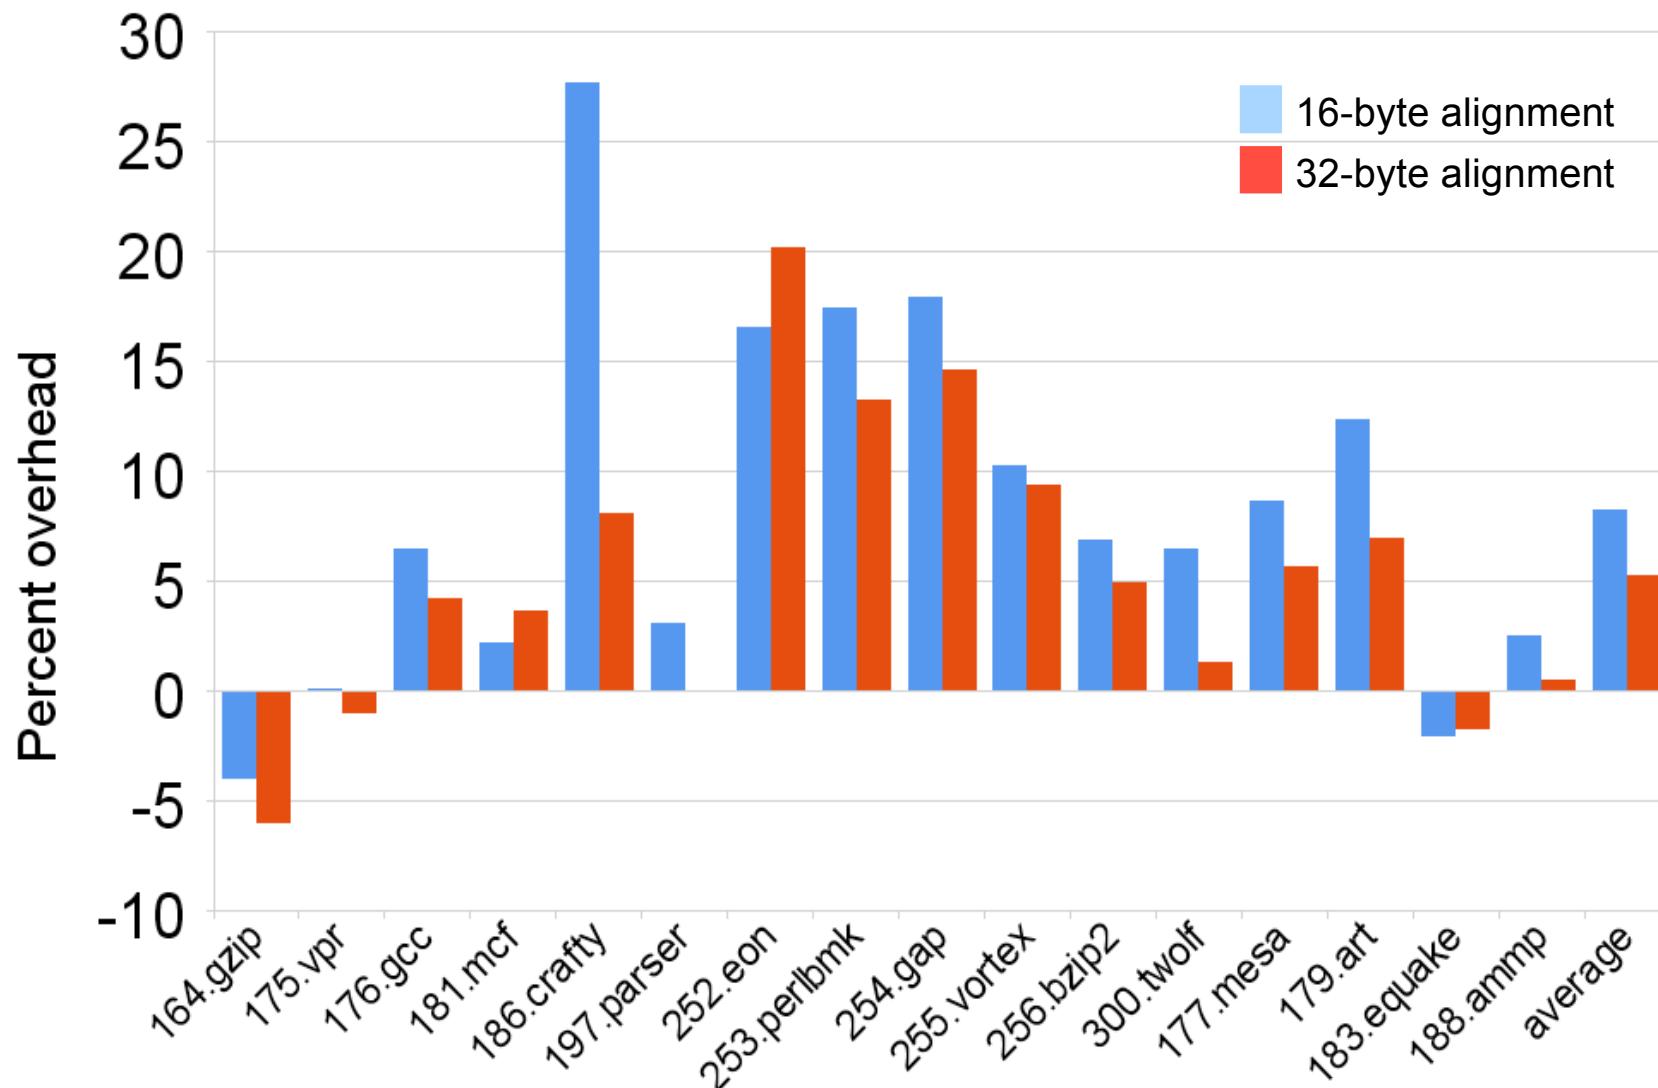
- inner sandbox
 - binary validation
 - hardware errata
- outer sandbox: OS system-call interception
- binary module loader
- trampoline interfaces
- IMC communications interface
- NPAPI interface
 - will likely be further constrained

Outline

- ◆ Introduction
- ◆ Background
- ◆ Eliminating Side Effects (SFI)
- ◆ Allowing Side-Effects Safely (Architecture)
- ◆ Evaluation

Performance

SPEC CPU2000 Benchmarks



What have we done with NaCl?

- Xaos
- Quake
- Lua interpreter
- H.264 decoder
- zlib (compression)
- cairo (2D rendering)
- bullet (physics)
- Life
- Voronoi diagram
- Spinning globe
- Mandelbrot
- ...

Related Work

There are many relevant comparisons:

- SFI
 - Wahbe, Lucco, Anderson, Graham
 - VINO (SFI), SPIN (type safety), Nooks (page)
 - MiSFIT, PittSFleld
- DTrace, Systemtap, XFI, Xax
- VMs: VMWare, Java
- JavaScript, Flash/AIR, Silverlight
- SSL, SSH, PGP

Questions?

<http://code.google.com/p/nativeclient>

Getting Involved

- Find a bug
- Port something
 - a compiler
 - an interpreted language
 - a game
 - an interesting library
- Build an application
- Build a multithreaded browser
- Add dynamic loading to NaCl
- Add support for JITted languages to NaCl